

Robotics / Intelligent Robotics

Introduction to Reinforcement Learning

**Luís Paulo Reis, Nuno Lau, Henrique L. Cardoso,
Armando Sousa**

lpreis@fe.up.pt

Director of LIACC – Artificial Intelligence and Computer Science Lab.


**Associate Professor at DEI/FEUP – Informatics Engineering Department, Faculty of Engineering
of the University of Porto, Portugal**

President of APPIA – Portuguese Association for Artificial Intelligence

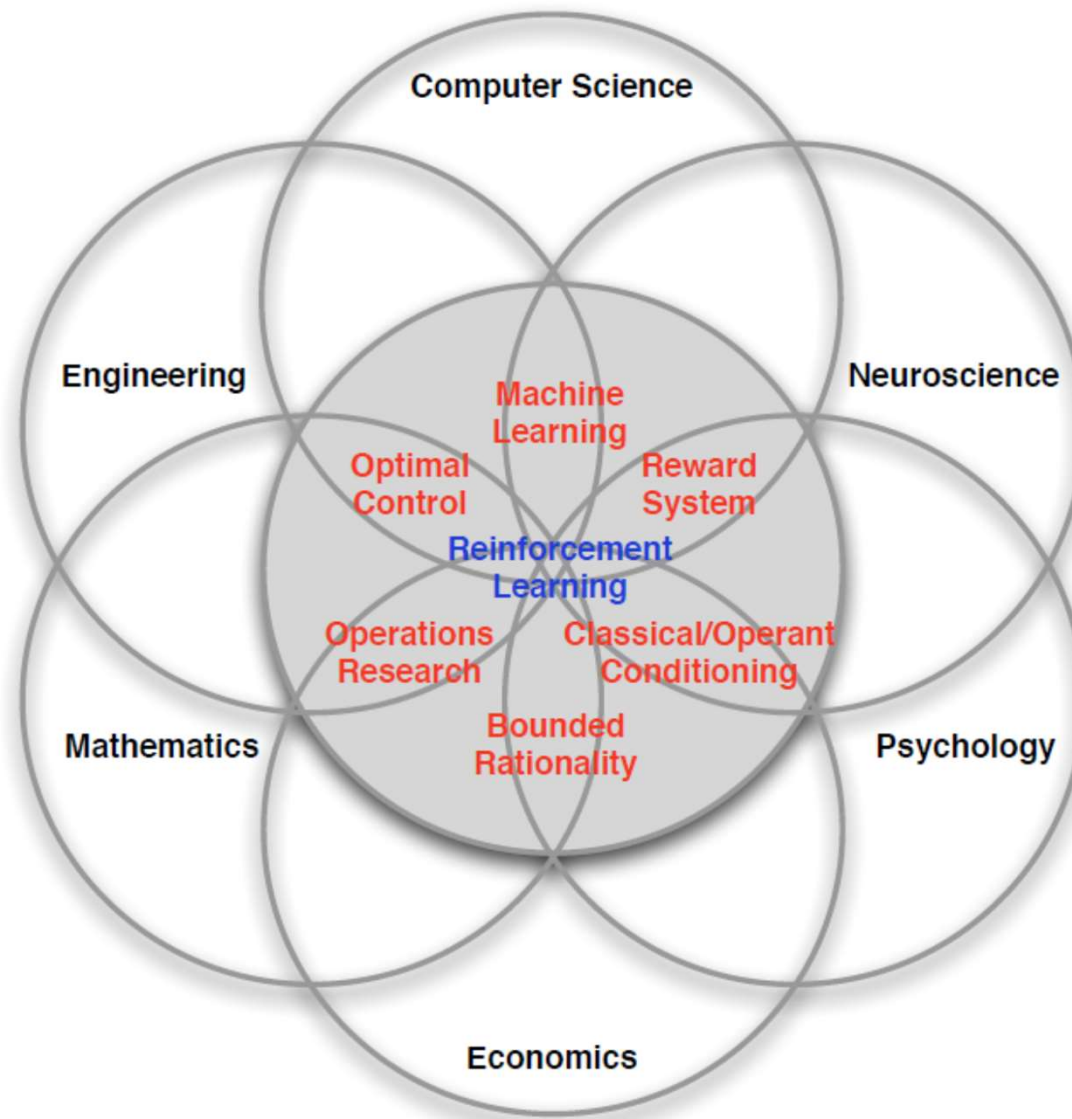


What is Reinforcement Learning?

- **Reinforcement Learning (RL)** is focused on **goal-directed learning from interaction**
- RL is **learning what to do** – how to map situations to actions – so as to maximize a numerical **return** signal
 - The learner is not told which actions to take: it must discover which actions yield the most **return** by trying them
 - **return** is the sum of rewards for a given sequence of actions
 - Typically, actions may affect not only immediate reward but also the next situation and subsequent rewards
- The **exploration-exploitation** tradeoff
 - Agent must prefer actions that it knows to be effective – **exploit**
 - But to discover such actions, it has to try actions not selected before – **explore**




What is Reinforcement Learning?



David Silver, RL slides

Branches of Machine Learning




David Silver, RL slides

RL vs (Un)Supervised Learning

- Different from **supervised learning**
 - In interactive problems it is impractical to obtain examples of desired behavior
 - In uncharted territory, an agent must learn from its own experience
- Different from **unsupervised learning**
 - RL is trying to maximize a reward signal, not trying to find hidden structure in collections of unlabeled data
- **RL** explicitly considers the *whole* problem of a **goal-directed agent interacting with an uncertain environment**
 - Creating a behavior model while applying it in the environment
- RL is the closest form of ML to the kind of learning humans do

Reinforcement Learning model




Learning to Play Tic-Tac-Toe

- Rule-based approach
 - Need to hardcode rules for each possible situations that might arise in a game
- Minimax
 - Assumes a particular way of playing by the opponent
- Dynamic programming can compute an optimal solution for any opponent
 - But requires as input a complete specification of that opponent (state/action probabilities)
- Can we obtain such information from **experience**?
 - Play many games against the opponent!

X	O	O
O	X	X
		X

Learning to Play Tic-Tac-Toe

- **States**
 - Possible configurations of the board
- **Actions**
 - Possible moves to make
- **Policy**
 - Which action should I play in each state?
- **Reward**
 - How good was the chosen action?



Elements of RL

- **Policy π**

- How should the agent behave over time?
- A policy is a (possibly stochastic) **mapping from perceived states to actions**

- **Reward signal r**

- Defines the **goal** of the RL problem
- On each time step, the environment sends a **reward** to the RL agent – the agent's goal is to **maximize the total reward (return)** received over the long run


- **Value function v**

- Specifies what is good in the long run
- The **value** of a state is the **total amount of reward** an agent can expect to accumulate from that state onwards (it takes into account future rewards)

→ We seek actions that bring about states of **highest value**, not highest reward, because these actions obtain the greatest **return**


Bandit Problems

- A simple setting with a **single state**



- K -armed bandit problem
 - There are k different **actions**
 - After each action a numerical reward is received from a stationary probability distribution
 - Each action has a **value** – its expected or mean reward, not known by the agent: $q_*(a) \doteq \mathbb{E}[R_t | A_t = a]$
 - The agent **estimates**, at time step t , the value of an action a : $Q_t(a)$

Bandit Problems



- Selecting *greedy* actions (whose estimated value is greatest): **exploiting**
- Selecting non-greedy actions: **exploring**
 - Improve estimates of non-greedy actions' value
- Lower reward in the short run (during *exploration*), but higher in the long run – after discovering the best actions, we can *exploit* them many times

Estimating Action Values

- Sample average:

$$Q_t(a) \doteq \frac{\sum_{i=1}^{t-1} R_i \cdot 1_{A_i=a}}{\sum_{i=1}^{t-1} 1_{A_i=a}}$$

- Update rule:

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

$$NewEstimate \leftarrow OldEstimate + StepSize [Target - OldEstimate]$$

- The target indicates a desirable direction in which to move
 - The *step-size parameter* changes from time step to time step
-
- Giving more weight to recent rewards – *constant step-size parameter*:

$$Q_{n+1} \doteq Q_n + \alpha [R_n - Q_n]$$

where $\alpha \in (0,1]$




Action Selection

- *Greedy* action selection (always exploits): $A_t \doteq \operatorname{argmax}_a Q_t(a)$
- ε -*greedy* action selection: behave greedily most of the time, but with small probability ε select randomly from among all the actions
 - $Q_t(a)$ will converge to $q_*(a)$ if a is selected sufficiently often
- *Soft-max* action selection (Boltzmann distribution):

$$\Pr\{A_t = a\} \doteq \frac{e^{Q_t(a)/\tau}}{\sum_{b=1}^k e^{Q_t(b)/\tau}}$$

where τ is a temperature parameter: if high, actions will tend to be equiprobable; if low, action values matter more; if $\tau \rightarrow 0$, then we have greedy action selection


The 10-armed Testbed



Markov Decision Processes

- In the general setting we have **many states**
- **Markov Decision Processes (MDP)** are a classical formalization of sequential decision making
 - Actions influence not just immediate rewards, but also subsequent situations (states) and thus future rewards
- A finite MDP, is defined by:
 - a set of states, S
 - a set of actions, $A(s)$
 - a state transition model, $p(s,a,s') \rightarrow [0,1]$
 - a reward function, $R(s,a,s') \rightarrow r$
- In MDPs we estimate the value $q_*(s, a)$

Agent-Environment Interface



- Dynamics of the MDP:

$$p(s', r | s, a) \doteq \Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\}$$

- The probability of each possible value for s' and r depends only in the immediately preceding state s and action a
- The state must include all relevant information about the past agent-environment interaction – **Markov property**

Example: Recycling Robot

- A robot has to decide whether it should (1) actively search for a can, (2) wait for someone to bring it a can, or (3) go to home base and recharge
- Searching is better (higher probability of getting a can) but runs down battery; if out of battery, the robot has to be rescued
- Decisions made on the basis of current energy level: high, low
- Reward is zero except when getting a can, and negative if out of battery

$$\mathcal{S} = \{high, low\}$$


$$\mathcal{A}(high) = \{search, wait\}$$

$$\mathcal{A}(low) = \{search, wait, recharge\}$$


$$r_{search} > r_{wait}$$

s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	-

Example: Recycling Robot




Example: maze



- Reward=-1 for each move
- Value function of optimal policy?

Example: maze



- Value function of optimal policy
- How to determine it?

Goals and Rewards

- A **reward signal** is used to define the **goal** of the agent
 - Learning to walk: reward proportional to the robot's forward motion
 - Learning to Escape from a maze: reward -1 for any state prior to escape (encourage escaping as quickly as possible)
 - Learning to find empty cans for recycling: reward of 0 most of the time, +1 for each can collected
 - Learning to play checkers or chess: reward +1 for winning, -1 for losing, and 0 for drawing and nonterminal positions
 - Provide **rewards** in such a way that by **maximizing** them the agent will also achieve our **goal**
 - The agent's goal is to **maximize the cumulative reward** it receives in the long run
- The reward signal is a way of communicating to the robot **what** you want it to achieve, not **how**

Returns and Episodes

- Agent wants to maximize **expected return**

$$G_t \doteq R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T$$


- **Episodic tasks**: when the agent-environment interaction breaks naturally into subsequences – **episodes**
 - From a **starting state** to a **terminal state**
 - Followed by a reset to another starting state, chosen independently of how the previous episode ended
- **Continuing tasks** do not break naturally into identifiable episodes (*e.g.*, on-going process-control)
 - Problem with calculating G_t :
 - $T = \infty$
 - G_t could also be infinite (if rewards are positive at each time step)

Returns and Episodes

- Adding **discounting**: agent wants to maximize **expected discounted return**


$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \qquad G_t \doteq R_{t+1} + \gamma G_{t+1}$$

- $0 \leq \gamma \leq 1$ is the **discount rate**
 - If $\gamma = 0$ the agent is “myopic” (only immediate reward matters)
 - As γ approaches 1, the agent becomes more farsighted (strongly considers future rewards)
- G_t is now finite, even if summing an infinite number of terms
- Applicable also to episodic tasks, if we consider a final absorbing state:



Example: Pole Balancing

- Move a **cart** so as to keep a **pole** from falling over
 - Failure if the pole falls past a given angle or if the cart runs off the track
 - The pole is reset to vertical after each failure



- **Episodic task**: reward +1 except when failure
 - return is the number of steps until failure
- **Continuing task**, using discounting: reward -1 on each failure and 0 otherwise
 - return is $-\gamma^K$, where K is the number of steps before failure


Policies and Value Functions

- RL algorithms involve estimating **value functions**
 - How good (in terms of expected return) is it to be in a given state?
 - How good is it to perform a given action in a given state?
- Future rewards depend on the choice of actions
 - Value functions are defined with respect to **policies** (ways of acting)
- **Policy**: a mapping from states to probabilities of selecting each possible action
 - $\pi(a|s) = \Pr(A_t = a|S_t = s)$

→ RL methods specify how the policy (*i.e.*, the probability distribution over $a \in \mathcal{A}(s)$ for each $s \in \mathcal{S}$) is changed with experience

Policies and Value Functions

- **State-value** function $v_{\pi}(s)$
 - Expected return when starting in s and following π thereafter
 - $v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t | S_t = s]$
- **Action-value** function $q_{\pi}(s, a)$
 - Expected return when taking action a in state s , and following π thereafter
 - $q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$
- The value functions v_{π} and q_{π} can be estimated from experience
 - $v_{\pi}(s) \doteq \sum_{s'} p(s, \pi(s), s') [r(s, \pi(s), s') + \gamma v_{\pi}(s')] \quad \text{(Bellman equation)}$
- Example: using a random policy, with $\gamma = 0.9$:



Gridworld



- Off-grid actions have no effect, with $r = -1$
- Any action from A gets to A', with $r = +10$
- Any action from B gets to B', with $r = +5$


3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

v_{π}

Optimal Policy and Value Function


- **Optimal policy π_*** : expected return is greater than any other policy

- **Optimal state-value function**: $v_*(s) \doteq \max_{\pi} v_{\pi}(s)$
- **Optimal action-value function**: $q_*(s, a) \doteq \max_{\pi} q_{\pi}(s, a)$




- Once we know v_* or q_* , the **optimal policy is greedy**

- Example:



Gridworld





22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

v_*

→	↕	←	↕	←
↶	↑	↷	←	←
↶	↑	↷	↷	↷
↶	↑	↷	↷	↷
↶	↑	↷	↷	↷

π_*

Example Grid World



- A bot is required to traverse a grid of 4×4 dimensions to reach its goal (1 or 16)
- There are 2 terminal states (1 and 16) and 14 non-terminal states (2 to 15)
- Each step is associated with a reward of -1
- Consider a random policy: at every state, the probability of every action {up, down, left, right} is 0.25
- Initialize v_1 for the random policy with all 0s

Example Grid World: Policy Evaluation

- Turning Bellman equation into an update:

$$v_{k+1}(s) \doteq \mathbb{E}_{\pi}[R_{t+1} + \gamma v_k(s_{t+1}) | S_t = s]$$

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

$$\begin{aligned}
 v_1(6) &= \sum_{a \in \{u,d,l,r\}} \pi(a|6) \sum_{s',r} p(s',r|6,a) [r + \gamma v_0(s')] \\
 &= \sum_{a \in \{u,d,l,r\}} \underbrace{\pi(a|6)}_{= 0.25 \forall a} \sum_{s'} \underbrace{p(s'|6,a)}_{= -1} \underbrace{[r + \gamma v_0(s')]}_{= 0 \forall s'} \\
 &= 0.25 * \{-p(2|6,u) - p(10|6,d) - p(5|6,l) - p(7|6,r)\} \\
 &= 0.25 * \{-1 - 1 - 1 - 1\} \\
 &= -1 \\
 &\Rightarrow v_1(6) = -1
 \end{aligned}$$

- For non-terminal states, $v_1(s) = -1$
- For terminal states, $p(s'|s, a) = 0$ (and hence $v_k(1) = v_k(6) = 0$, for all k)

$v_1 =$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

Example Grid World: Policy Evaluation

- Step 2, with discount factor $\gamma = 1$

$$\begin{aligned}
 v_2(6) &= \sum_{a \in \{u,d,l,r\}} \underbrace{\pi(a|6)}_{= 0.25 \forall a} \sum_{s'} p(s'|6,a) \underbrace{[r + \gamma v_1(s')]}_{= -1} = \begin{cases} -1, s' \in S \\ 0, s' \in S^+ \setminus S \end{cases} \\
 &= 0.25 * \{p(2|6,u)[-1 - \gamma] + p(10|6,d)[-1 - \gamma] \\
 &\quad + p(5|6,l)[-1 - \gamma] + p(7|6,r)[-1 - \gamma]\} \\
 &\stackrel{\gamma=1}{=} 0.25 * \{-2 - 2 - 2 - 2\} \\
 &= -2
 \end{aligned}$$

- For all red states, $v_2(s) = -2$

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16


- For the other states (2, 5, 12, 15):

$$\begin{aligned}
 v_2(2) &= \sum_{a \in \{u,d,l,r\}} \underbrace{\pi(a|2)}_{= 0.25 \forall a} \sum_{s'} p(s'|2,a) \underbrace{[r + \gamma v_1(s')]}_{= -1} = \begin{cases} -1, s' \in S \\ 0, s' \in S^+ \setminus S \end{cases} \\
 &= 0.25 * \{p(2|2,u)[-1 - \gamma] + p(6|2,d)[-1 - \gamma] \\
 &\quad + p(1|2,l)[-1 - \gamma * 0] + p(3|2,r)[-1 - \gamma]\} \\
 &\stackrel{\gamma=1}{=} 0.25 * \{-2 - 2 - 1 - 2\} \\
 &= -1.75 \\
 &\Rightarrow v_2(2) = -1.75
 \end{aligned}$$

$v_2 =$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

Example Grid World: Policy Evaluation



Random policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0


$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

$\leftarrow v_{\pi}$



Optimal policy

Policy Evaluation

Algorithm 1 Policy Iteration (Policy Evaluation)

Require: $V(s)$ arbitrarily initialized, $\forall s \in S$

```
1: repeat
2:   repeat
3:      $\Delta \leftarrow 0$ 
4:     for all  $s \in S$  do
5:        $v \leftarrow V(s)$ 
6:        $V(s) \leftarrow \sum_{s'} P(s, \pi(s), s') [r(s, \pi(s), s') + \gamma V(s')]$ 
7:        $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
8:     end for
9:   until  $\Delta \leq \theta$  ▷ a small positive value
10:  ...
11: until StablePolicy == True
```



Policy Improvement

Algorithm 2 Policy Iteration (Policy Improvement)

Require: $V(s)$ arbitrarily initialized, $\forall s \in S$

```
1: repeat
2:   ...
3:   StablePolicy  $\leftarrow$  True
4:   for all  $s \in S$  do
5:      $b \leftarrow \pi(s)$ 
6:      $\pi(s) \leftarrow \arg \max_a \sum_{s'} P(s, a, s')[r(s, a, s') + \gamma V(s')]$ 
7:     if  $b \neq \pi(s)$  then
8:       StablePolicy  $\leftarrow$  False
9:     end if
10:  end for
11: until StablePolicy == True
```



Value Iteration

Algorithm 3 Value Iteration

Require: $V(s)$ arbitrarily initialized, $\forall s \in S$


```
1: repeat  
2:    $\Delta \leftarrow 0$   
3:   for all  $s \in S$  do  
4:      $v \leftarrow V(s)$   
5:      $V(s) \leftarrow \max_a \sum_{s'} P(s, a, s') [r(s, a, s') + \gamma V(s')]$   
6:      $\Delta \leftarrow \max(\Delta, \|v - V(s)\|)$   
7:   end for  
8: until  $\Delta \leq \theta$  ▷ a small positive value
```

Optimal Policy from Value Function

- How to determine the policy from the Value Function?
- Extract a policy!
- Policy maps states to actions

$$\pi(s) = \arg \max_a \sum_{s'} P(s, a, s') [r(s, a, s') + \gamma V(s')]$$

Example: Maze Optimal Policy



Approximation

- Optimal policies are **computationally costly** to find – we can only approximate
 - In tasks with small, finite state sets: **tabular methods**
 - Otherwise: function approximation using a more compact parameterized function representation (e.g. using neural networks)

→ The online nature of RL allows us to *put more effort into learning to make decisions for frequently encountered states*

Temporal-Difference Learning

- TD methods update estimates based on immediately observed reward and state

- Update rule:

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

- Because TD bases its update in part on an existing estimate, it is a *bootstrapping* method

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$


$S \leftarrow S'$

 until S is terminal

Temporal-Difference Learning

- TD vs Dynamic Programming methods
 - TD methods **do not require a model** of the environment's dynamics (rewards and next-state probability distributions)
- TD vs Monte Carlo methods
 - TD methods are naturally implemented in an **online, fully incremental fashion**, while MC methods must wait until the end of an episode
 - Useful if episodes are very long, or in continuing tasks (that have no episodes at all)
- Usually, TD methods converge faster than MC methods on stochastic tasks

Sarsa: On-policy TD Control



- Update rule:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

- This rule uses every element of the quintuple of events $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1})$

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ε -greedy)

Loop for each step of episode:


Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal



- Converges to optimal policy and action-value function if all state-action pairs are visited infinitely and policy converges to greedy (e.g. using ε -greedy with $\varepsilon = 1/t$)

Q-learning: Off-policy TD Control

- Update rule:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:


 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

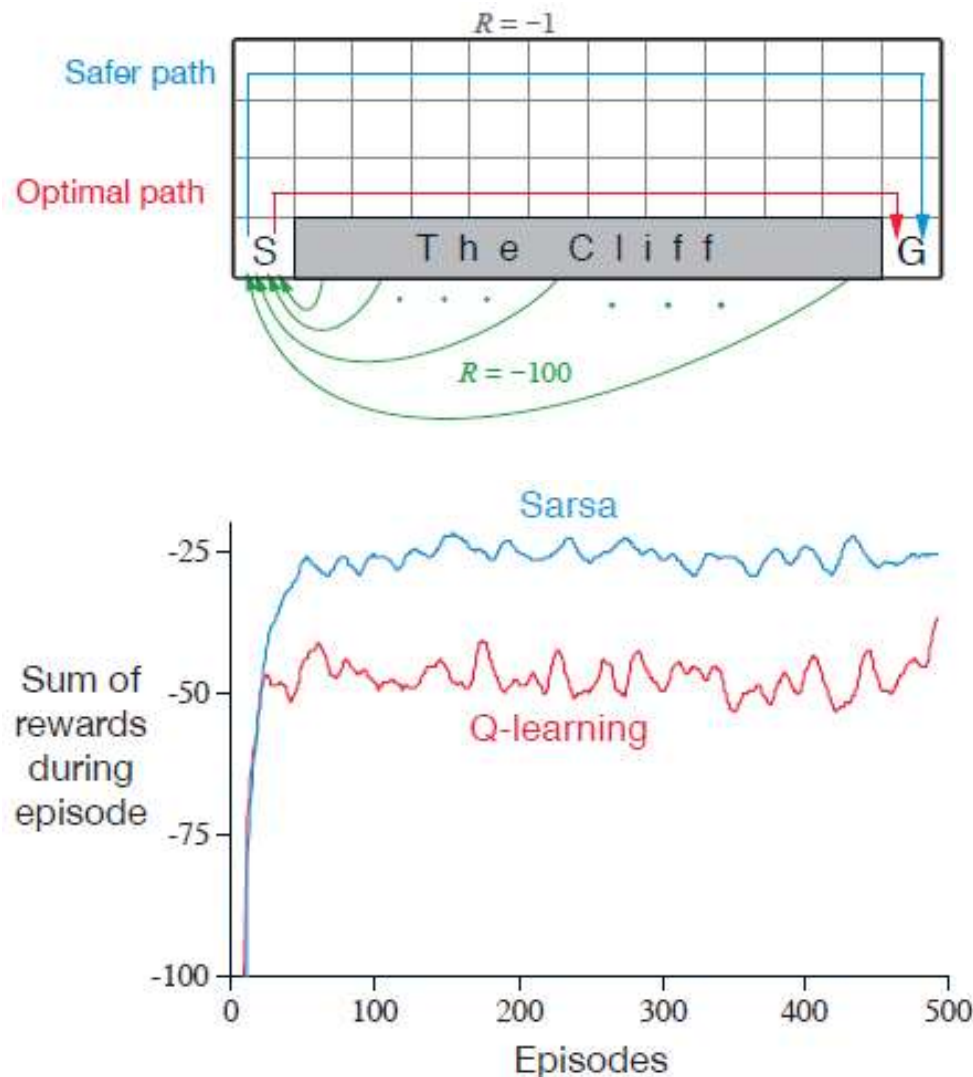
$S \leftarrow S'$

 until S is terminal




- The learned action-value function Q directly approximates q_* , independently of the policy being followed

Example: Cliff Walking



- Sarsa and Q-learning with ϵ -greedy action selection ($\epsilon = 0.1$)
 - Q-learning learns values for the optimal policy
 - Sarsa takes action selection into account and learns the longer but safer path
 - Given exploration, Q-learning occasionally falls off the cliff, hence the lower online performance
- If ϵ is gradually reduced, both methods converge to the optimal policy

Example: Maze




Classes of Reinforcement Learning

Three **main RL classes** of methods

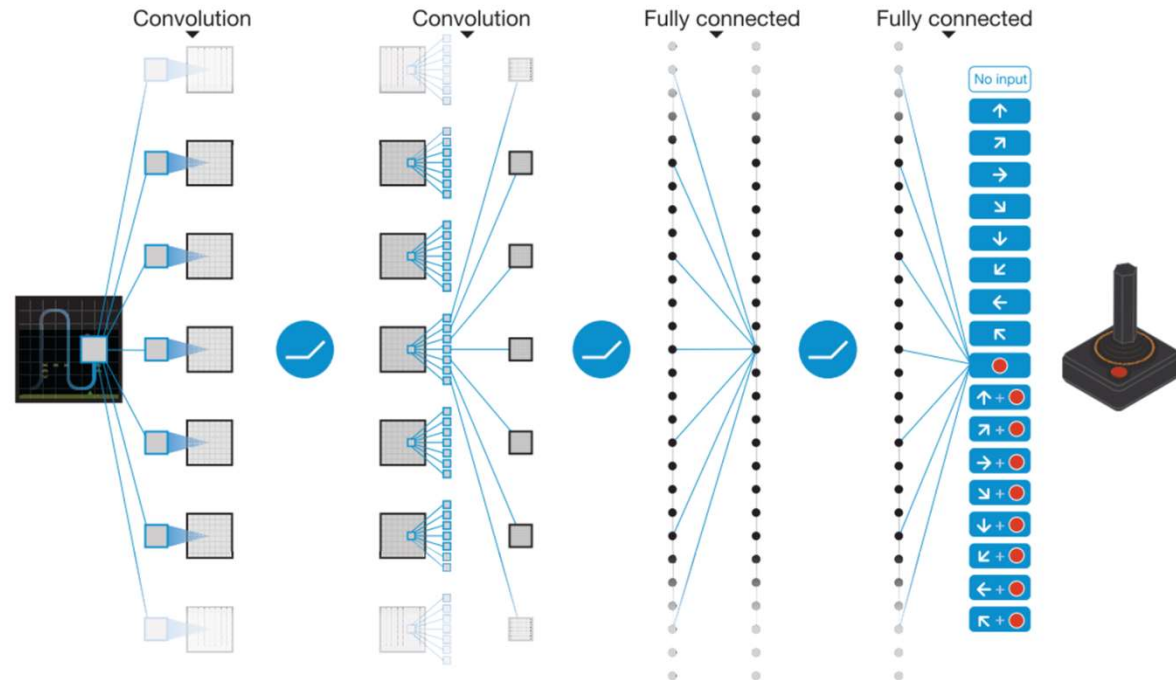
- **Value Function based** methods
 - No policy representation
 - Policy obtained by evaluating the value function directly
- **Policy Search** methods
 - No value function
 - Optimization of a parametrized policy directly on policy-space
- **Actor-Critic** methods
 - Value function (critic)
 - Explicit Policy representation (actor)
- **Batch RL is a sub-class of Value Function based methods**

Batch Reinforcement Learning

- **Batch RL** estimates value functions by processing a **set of interactions**
- The value function is updated synchronously
- Application of function approximators
- Collected experience is **not discarded**
- Data **efficient**
- Fit $\bar{Q}_i = r_i + \gamma \max_b \hat{Q}(s_{i+1}, b)$

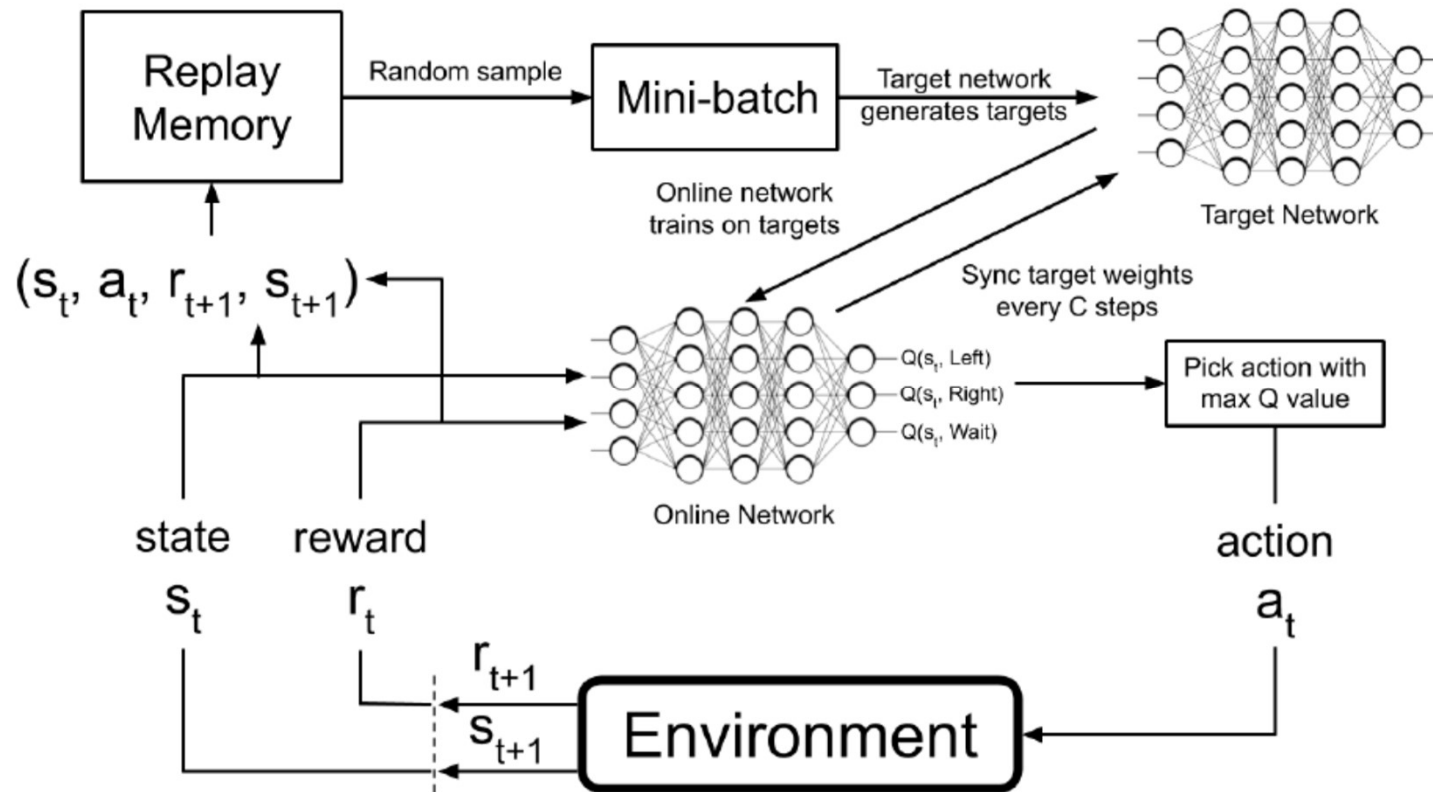


Deep Q Network - DQN




- Input: raw pixels (84x84x4 image)
- Output: Q values for all possible actions

Deep Q Network - DQN



$$\bar{Q}_i = r_i + \gamma \max_b \hat{Q}(s_{i+1}, b)$$

Deep Q Network - DQN



Open AI Gym

```
import gym

# load the environment
env = gym.make('Example')

# perform N episodes
while True:

    # prepare the environment for the next episode
    obs = env.reset()

    # flag for episode completion
    done = False

    # run the episode
    while not done:

        # choose how to act based on the current state
        # usually the output of a neural network
        action = env.action_space.sample()

        # advance the simulation by one timestep
        # by interacting with the world
        obs, reward, done, info = env.step(action)


# cleanup
env.close()
```

RL Algorithms

Algorithm	Description	Policy	Action Space	State Space
Monte Carlo	Every visit to Monte Carlo	Either	Discrete	Discrete
Q-learning	State–action–reward–state	Off-policy	Discrete	Discrete
SARSA	State–action–reward–state–action	On-policy	Discrete	Discrete
Q-learning - Lambda	Q-learning with eligibility traces	Off-policy	Discrete	Discrete
SARSA - Lambda	SARSA with eligibility traces	On-policy	Discrete	Discrete
DQN [Mnih et al., 2013]	Deep Q Network	Off-policy	Discrete	Continuous
DDPG [Lillicrap et al., 2016]	Deep Deterministic Policy Gradient	Off-policy	Continuous	Continuous
A3C [Mnih et al., 2016]	Asynchronous Advantage Actor-Critic	On-policy	Continuous	Continuous
NAF [Gu et al., 2016]	Q-Learning with Normalized Advantage Functions	Off-policy	Continuous	Continuous
TRPO [Schulman et al., 2015]	Trust Region Policy Optimization	On-policy	Continuous	Continuous
PPO [Schulman et al., 2017]	Proximal Policy Optimization	On-policy	Continuous	Continuous
TD3 [Fujimoto et al., 2018]	Twin Delayed Deep Deterministic Policy Gradient	Off-policy	Continuous	Continuous
SAC [Haarnoja et al., 2018]	Soft Actor-Critic	Off-policy	Continuous	Continuous

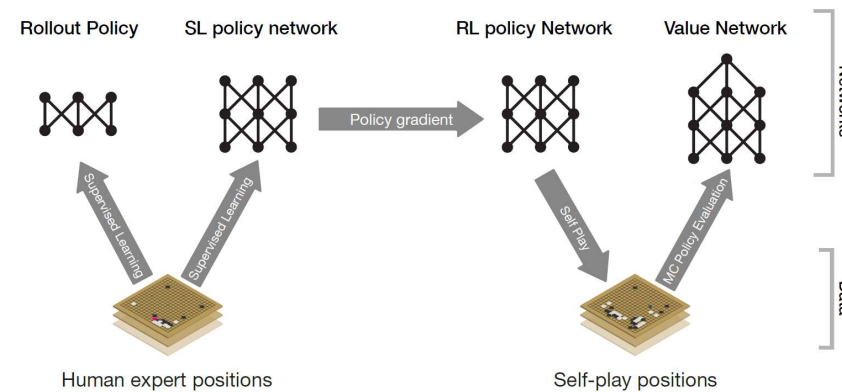

RL in Games

- TD-Gammon [Tesauro, 1995]
 - Neural Network trained with self-play reinforcement learning
- Atari 2600 Games [DeepMind, 2013]
 - Learn control policies directly from high-dimensional sensory input using reinforcement learning
 - Input is raw pixels and output is a value function estimating future rewards




RL in Games

- AlphaGo [Google DeepMind, 2016]
 - Convolutional Neural Networks trained with human expert data
 - Deep Reinforcement Learning with fictitious self-play

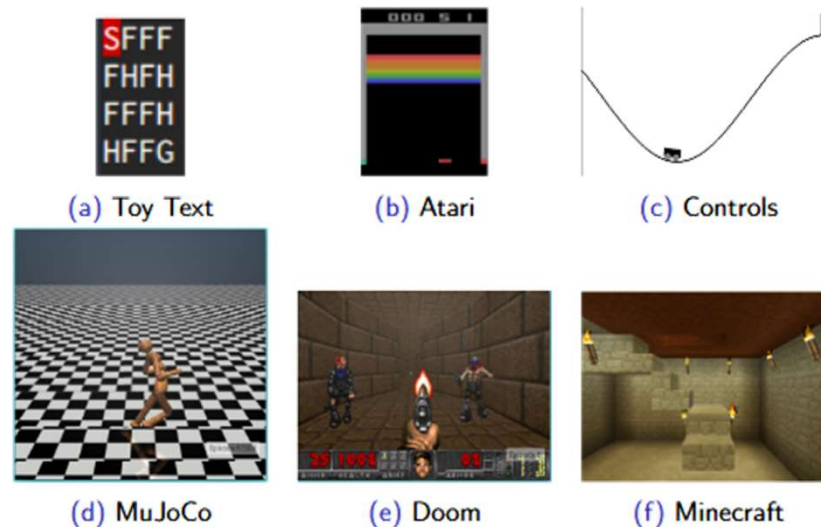


- Poker: Heads-Up Limit Texas Hold'em – NFSP [UCL, 2016]
 - Deep Reinforcement Learning with fictitious self-play
 - No prior knowledge



OpenAI Gym

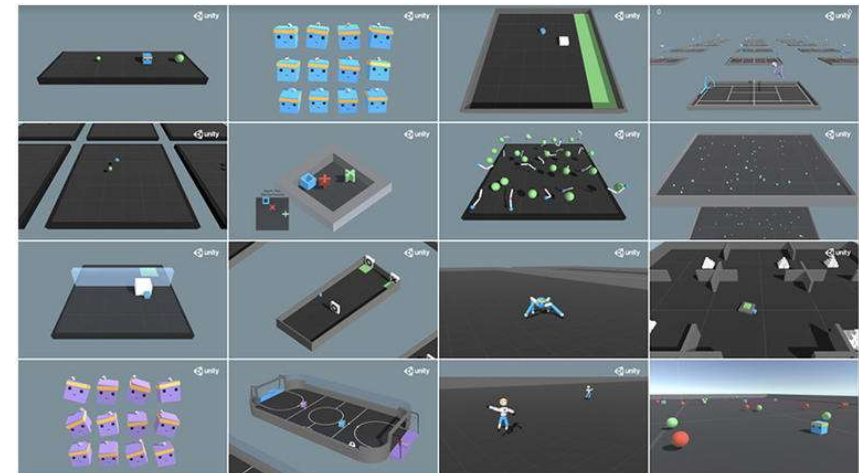
- [Gym](#) is a toolkit for developing and comparing reinforcement learning algorithms
- The [gym library](#) is a collection of test problems with a shared interface — **environments** — that you can use to work out your RL algorithms



- [OpenAI Baselines](#) is a set of high-quality implementations of RL algorithms
 - See also [Stable Baselines](#)

Unity ML-Agents

- With Unity Machine Learning Agents ([ML-Agents](#)), you teach intelligent agents through a combination of **deep reinforcement learning** and **imitation learning**



Further Reading

- Sutton, R. S. and Barto, A. G. (2018). *Reinforcement Learning – An Introduction*, 2nd ed., The MIT Press: Chap. 1-3, 6
- Simple Tutorial Videos for Deep RL:
 - Introduction on Reinforcement Learning and Deep RL:
<https://www.youtube.com/watch?v=JgvyzlkgxF0>
 - PPO – Proximal Policy Optimization (PPO):
<https://www.youtube.com/watch?v=5P7l-xPq8u8>

Conclusions

- RL enables to learn intelligent behavior in complex environments
- Large number of algorithms and approaches
- Amazing results in vintage Atari Games
- Stunning results of AlphaGo and AlphaZero
- Very promising results in Robotics
- Very fast evolution in the last few years

Robotics / Intelligent Robotics

Introduction to Reinforcement Learning

**Luís Paulo Reis, Nuno Lau, Henrique L. Cardoso,
Armando Sousa**

lpreis@fe.up.pt

Director of LIACC – Artificial Intelligence and Computer Science Lab.

**Associate Professor at DEI/FEUP – Informatics Engineering Department, Faculty of Engineering
of the University of Porto, Portugal**

President of APPIA – Portuguese Association for Artificial Intelligence

